



SLURM @UPPMAX

Diana Iușan

*Application expert and
training coordinator UPPMAX*



UPPSALA
UNIVERSITET

More Slurm and other advanced UPPMAX techniques

- A closer look at Slurm
- Using the GPUs on Snowy
- Job efficiency with the **jobstats** tool
- Advanced job submission





From login to the first job script



Technical summary of the UPPMAX clusters

	Rackham	Snowy	Bianca
Purpose	General-purpose	General-purpose	Sensitive
# Intel CPU Nodes	486	228	288
# GPU Nodes	-	50, Nvidia T4	10, 2x Nvidia A100 each
Cores per node	20	16	16 or 64
Memory per node	128 GB	128 GB	128 GB
Fat nodes	256 GB & 1 TB	256, 512 GB & 4 TB	256 & 512 GB
Local disk (scratch)	2/3 TB	4 TB	4 TB
Login nodes	Yes (4)	No (reached from Rackham)	Yes (2 cores and 15 GB)
"Home" storage	Domus	Domus	Castor
"Project" Storage	Crex, Lutra	Crex, Lutra	Castor



Technical summary of the UPPMAX clusters

	Rackham	Snowy	Bianca
Purpose	General-purpose	General-purpose	Sensitive
# Intel CPU Nodes	486	228	288
# GPU Nodes	-	50, Nvidia T4	10, 2x Nvidia A100 each
Cores per node	20	16	16 or 64
Memory per node	128 GB	128 GB	128 GB
Fat nodes	256 GB & 1 TB	256, 512 GB & 4 TB	256 & 512 GB
Local disk (scratch)	2/3 TB	4 TB	4 TB
Login nodes	Yes (4)	No (reached from Rackham)	Yes (2 cores and 15 GB)
"Home" storage	Domus	Domus	Castor
"Project" Storage	Crex, Lutra	Crex, Lutra	Castor



The **Slurm** Workload Manager

- provides a framework for starting, executing, and monitoring jobs on the compute nodes



The **Slurm** Workload Manager

- provides a framework for starting, executing, and monitoring jobs on the compute nodes
- schedules the jobs on the clusters
- allocates the required resources (compute cores or nodes, memory)



The **Slurm** Workload Manager

- Free, popular, lightweight
- Open source:
<https://slurm.schedmd.com>
- available at all SNIC centra
- UPPMAX Slurm userguide:
https://docs.uppmax.uu.se/cluster_guides/slurm/





Running on the login node vs. job submission



How to submit a job?



- Recap:

sbatch	-A uppmax2025-2-262	-t 10:00	-p core	-n 10	my_job.sh
Slurm batch	Project name	Maximum runtime	Partition (“job type”)	#cores	job script



Job time limits

- `-t` minutes
- `-t` minutes:seconds
- `-t` hours:minutes:seconds
- `-t` days-hours
- `-t` days-hours:minutes
- **`-t days-hours:minutes:seconds`**





What should the wall time of a job be?

- Your first job? Testing new software or input?





What should the wall time of a job be?

- Your first job? Testing new software or input?
 - Use a short time limit, 10 min - 1h.





What should the wall time of a job be?

- Your first job? Testing new software or input?
 - Use a short time limit, 10 min - 1h.
- Q: When you have no idea how long a program will take to run, what should you book?
- Q: When you have an idea of how long a program would take to run, what should you book?



What should the wall time of a job be?

- Your first job? Testing new software or input?
 - Use a short time limit, 10 min - 1h.
- Q: When you have no idea how long a program will take to run, what should you book?
 - A: very long time, e.g. 10-00:00:00
- Q: When you have an idea of how long a program would take to run, what should you book?
 - A: overbook by 50%



Technical summary of the UPPMAX clusters

	Rackham	Snowy	Bianca
Purpose	General-purpose	General-purpose	Sensitive
# Intel CPU Nodes	486	228	288
# GPU Nodes	-	50, Nvidia T4	10, 2x Nvidia A100 each
Cores per node	20	16	16 or 64
Memory per node	128 GB	128 GB	128 GB
Fat nodes	256 GB & 1 TB	256, 512 GB & 4 TB	256 & 512 GB
Local disk (scratch)	2/3 TB	4 TB	4 TB
Login nodes	Yes (4)	No (reached from Rackham)	Yes (2 cores and 15 GB)
"Home" storage	Domus	Domus	Castor
"Project" Storage	Crex, Lutra	Crex, Lutra	Castor



Technical summary of the UPPMAX clusters

	Rackham	Snowy	Bianca
Purpose	General-purpose	General-purpose	Sensitive
# Intel CPU Nodes	486	228	288
# GPU Nodes	-	50, Nvidia T4	10, 2x Nvidia A100 each
Cores per node	20	16	16 or 64
Memory per node	128 GB	128 GB	128 GB
Fat nodes	256 GB & 1 TB	256, 512 GB & 4 TB	256 & 512 GB
Local disk (scratch)	2/3 TB	4 TB	4 TB
Login nodes	Yes (4)	No (reached from Rackham)	Yes (2 cores and 15 GB)
"Home" storage	Domus	Domus	Castor
"Project" Storage	Crex, Lutra	Crex, Lutra	Castor



More on partitions

- **-p core**
 - “core” is the default partition
 - ≤ 20 cores on Rackham
 - ≤ 20 cores on Bianca
 - a script or program written without any thought on parallelism will use 1 core



More on partitions



- **-p core**
 - “core” is the default partition
 - ≤ 20 cores on Rackham
 - ≤ 20 cores on Bianca
 - a script or program written without any thought on parallelism will use 1 core
- **-p node**
 - if you wish to book full node(s)





Getting on the fast lane

- for quick tests and development work



Quick testing



- The “devel” partition
 - max 2 nodes per job
 - up to 1 hour in length
 - only 1 at a time
 - **-p devcore, -p devel**



Quick testing



- The “devel” partition
 - max 2 nodes per job
 - up to 1 hour in length
 - only 1 at a time
 - **-p devcore, -p devel**
 - Any free nodes in the devel partition? Check status with
 - `sinfo -p devel`
 - `jobinfo -p devel`



Quick testing



- Any free nodes in the devel partition? Check status with
 - `sinfo -p devel`
 - **`jobinfo -p devel`**
 - more on these tools later
- High priority queue for short jobs
 - 4 nodes
 - up to 15 minutes
 - **--qos=short**





- What if 1h is not enough?
- What if the application uses a graphical user interface (GUI)?





Debugging or complicated workflows

- Interactive jobs
 - handy for debugging a code or a script by executing it line by line or for using programs with a graphical user interface
 - `salloc -n 80 -t 03:00:00 -A uppmax2025-2-262`
 - `interactive -n 80 -t 03:00:00 -A uppmax2025-2-262`





Debugging or complicated workflows

- Interactive jobs
 - handy for debugging a code or a script by executing it line by line or for using programs with a graphical user interface
 - `salloc -n 80 -t 03:00:00 -A uppmax2025-2-262`
 - `interactive -n 80 -t 03:00:00 -A uppmax2025-2-262`
 - up to 12 hours
 - useful together with the `--begin=<time>` flag
 - `salloc -A uppmax2025-2-262 --begin=2025-02-13T08:00:00` asks for an interactive job that will start earliest tomorrow at 08:00



Parameters in the job script or the command line?

- Command line parameters override script parameters
- A typical script may be:

```
#!/bin/bash

#SBATCH -A uppmax2025-2-262

#SBATCH -p core

#SBATCH -n 1

#SBATCH -t 24:00:00
```

- Just a quick test:

```
sbatch -p devcore -t 00:15:00 jobsript.sh
```



Hands-on #1:

sbatch/jobinfo

- login to Rackham
- find out which projects you're a member of using projinfo
- submit a short (10 min) test job; note the job ID
- find out if there are any free nodes in the devel partition
- submit a new job to use the devel partition
- write in the HackMD when you're done



Memory in core or devcore jobs

- `-n X`
- on Rackham: 6.4 GB per core
- on Snowy/Bianca: 8GB per core
- Slurm reports the available memory in the prompt at the start of an interactive job



More flags

- `-J <jobname>`
- email:
 - `--mail-type=BEGIN,END,FAIL,TIME_LIMIT_80`
 - `--mail-user` - Don't use. Set your email correctly in SUPR instead.
- out/err redirection:
 - `--output=slurm-%j.out` and `--error=slurm-%j.err`
by default, where `%j` will be replaced by the job ID
 - `--output=my.output.file`
 - `--error=my.error.file`



More flags

- Memory
 - -C thin / -C 128GB
 - -C fat / -C 256GB / -C 1TB
- Dependencies: --dependency
- Job array: --array
- More at <https://slurm.schedmd.com/sbatch.html>
 - or just man sbatch
 - not all options work on all systems!





From job submission to job monitoring

- sbatch
- sinfo
- jobinfo (wrapper around squeue)



Monitoring jobs



- **jobinfo** - a wrapper around **squeue**
 - lists running and pending jobs
 - `jobinfo -u username`
 - `jobinfo -A uppmax2025-2-262`
 - `jobinfo -u username --state=running`
 - `jobinfo -u username --state=pending`
- One may also use the **squeue** command.



Monitoring and modifying jobs

- **scontrol**
 - scontrol show job *jobid*





Monitoring and modifying jobs

- **scontrol**
 - `scontrol show job jobid`
- possible to modify the job details after the job has been submitted; some options, like maximum runtime, may be modified (=shortened) even after the job started
 - `scontrol update JobID=jobid QOS=short`
 - `scontrol update JobID=jobid TimeLimit=1-00:00:00`
 - `scontrol update JobID=jobid NumNodes=10`
 - `scontrol update JobID=jobid Features=mem1TB`



When a job goes wrong



- **scancel**
 - *jobid*
 - -u *username* - to cancel all your jobs
 - -t *state* - cancel pending or running jobs
 - -n *name* - cancel jobs with a given name
 - -i - ask for confirmation



Priority



- Roughly:
- The first job of the day has elevated priority
- Other normal jobs run in the order of submission (subject to scheduling)
- Projects exceeding their allocation get successively into the lower priority category
- Bonus jobs run after the jobs in the higher priority categories



Priority



- In practice:
 - submit early = run early
 - bonus jobs always run eventually, but may need to wait until the night or weekend
- In detail:

<https://docs.uppmax.uu.se/software/jobinfo/#initial-priority-at-submit-time>



Hands-on #2:

sbatch/squeue/scancel/scontrol/jobinfo

- submit a new job; note the job ID
- check all your running jobs
- what is the priority or your recently-submitted job?
- submit a new job to run for 24h; note the job ID
- modify the name of the job to “wrongjob” and the maximum runtime to 7days, for example
- cancel your job with name “wrongjob”





Determining job efficiency

- **jostats** - custom-made UPPMAX tool



Job efficiency



- **jobstats** - a tool in the fight for productivity
 - it works only for jobs longer than 5-15 minutes
 - **-r jobid** - check running jobs
 - **-A project** - check all recent jobs of a given project
 - **-p jobid** - produce a CPU and memory usage plot
 - **-M cluster** - check jobs on other cluster
 - <https://docs.uppmax.uu.se/software/jobstats/>





Hands-on #3: jobstats

- Generate jobstats plots for your jobs
 - Firstly, find some job IDs from this month
 - `finishedjobinfo -m username`
 - Write down the IDs from some interesting jobs.
 - Generate the images:
`$ jobstats -p ID1 ID2 ID3`
- Look at the images. You may find some interesting ones in
`/proj/introtouppmax/labs/moreslurm/jobstatsplots/`
`$ eog *png &`



Hands-on #3: jobstats

- Which of the plots in
`/proj/introtouppmax/labs/moreslurm/
jobstatsplots/`
 - Show good CPU or memory usage?
 - Indicate that the job requires a fat node?



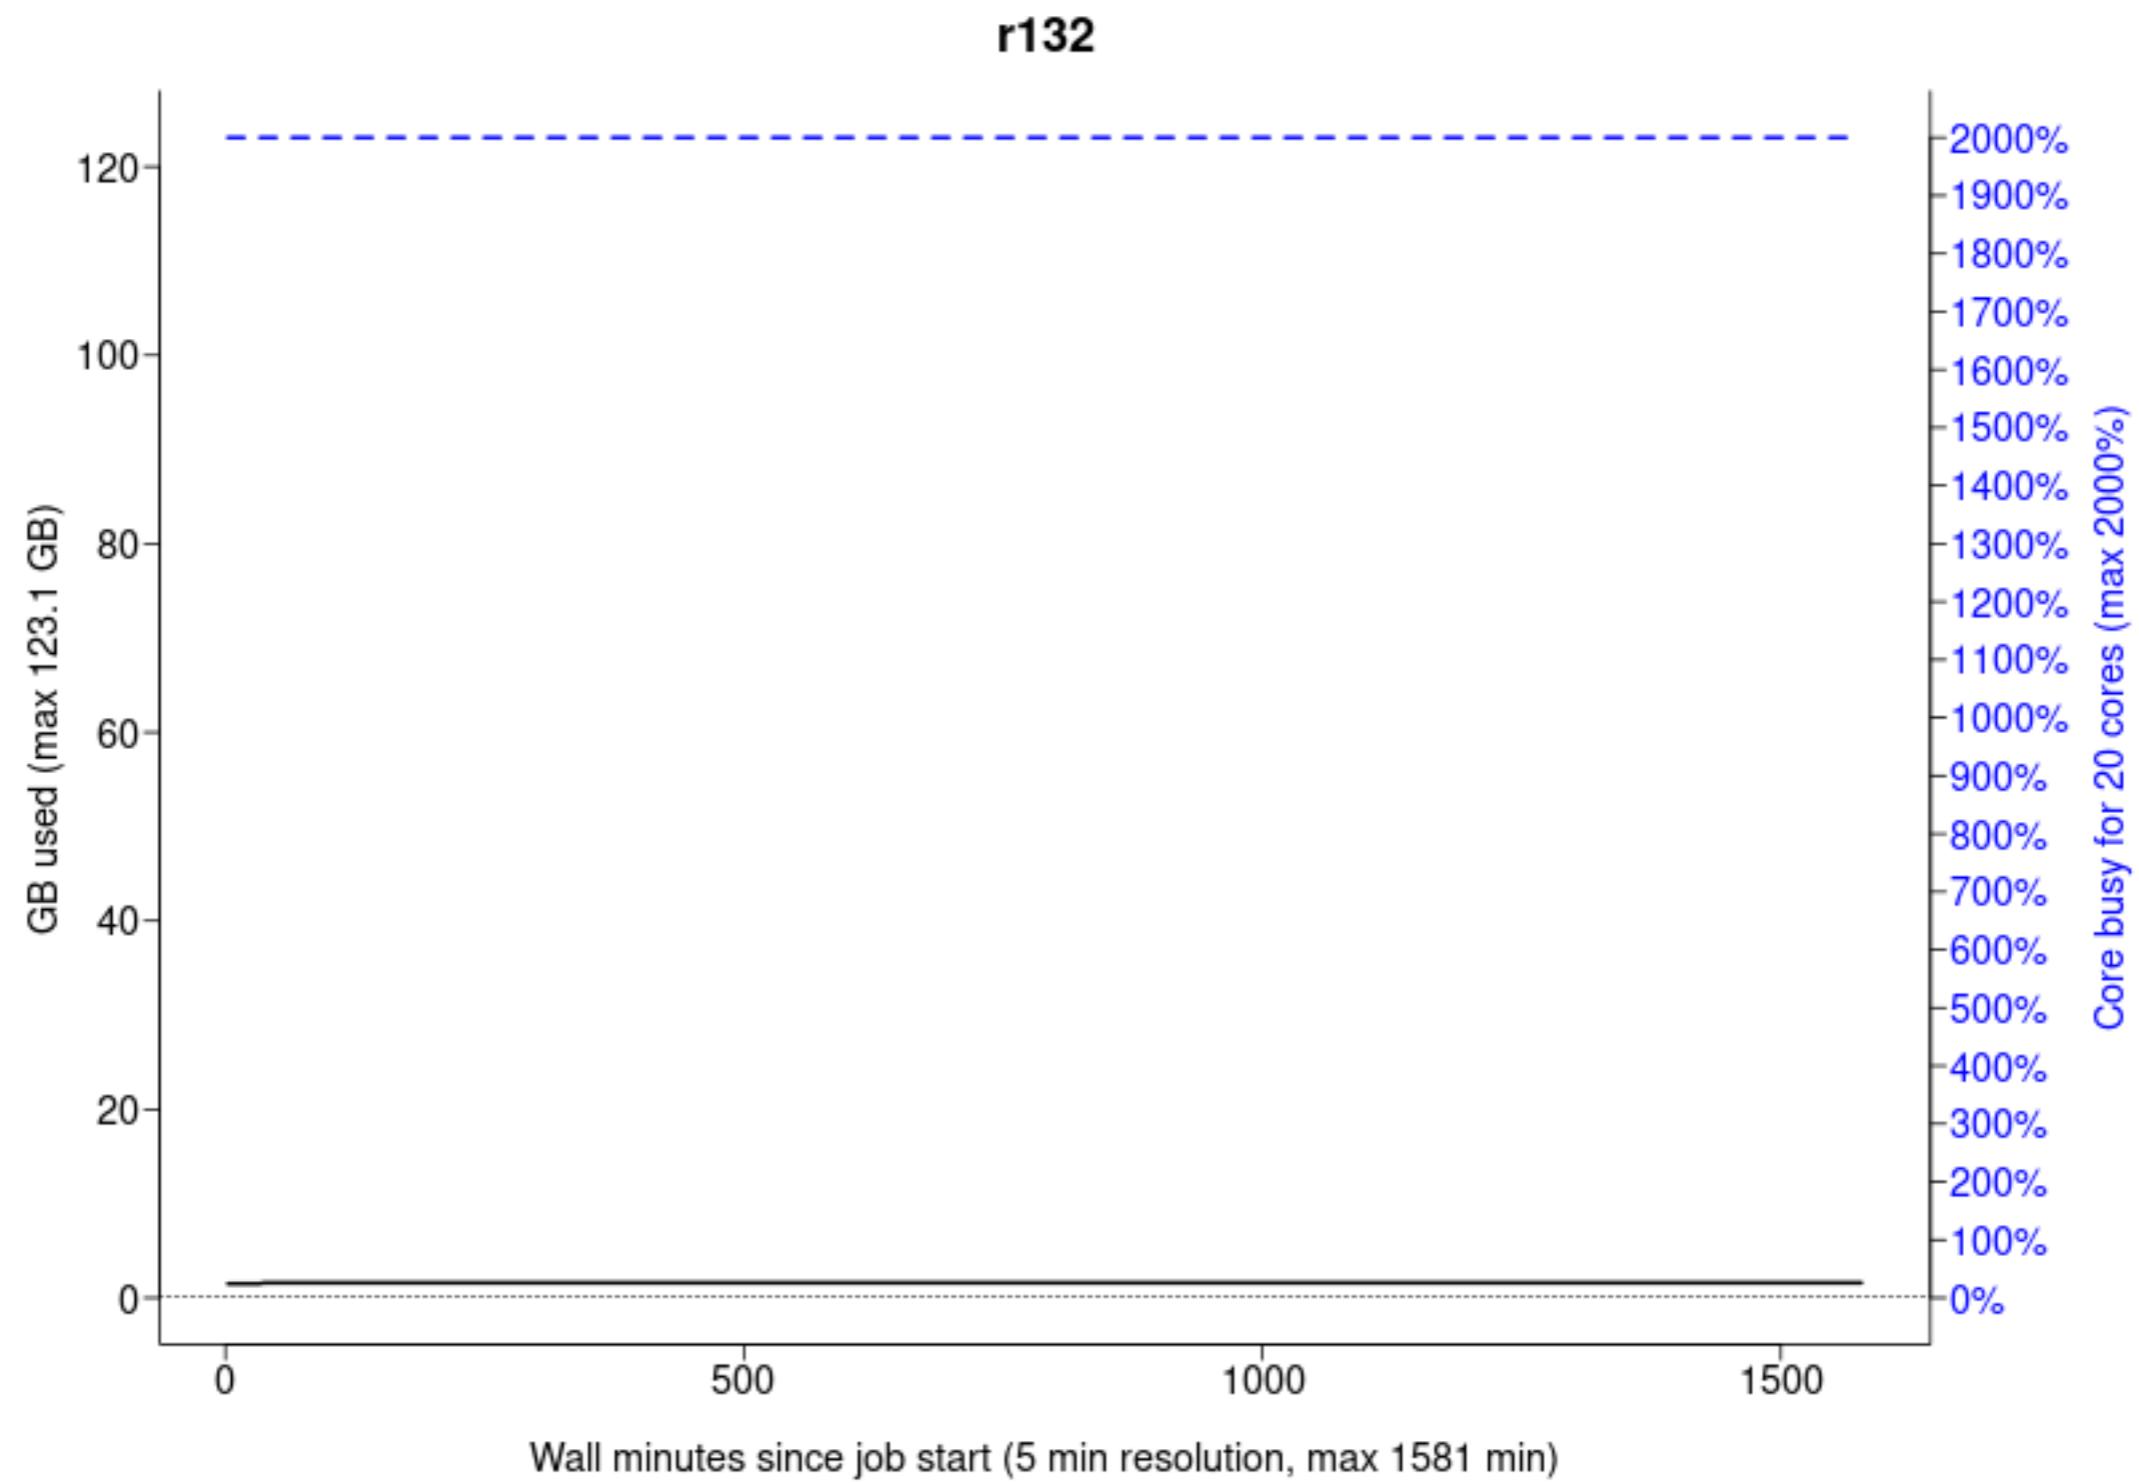
11217530 COMPLETED on rackham end: 2019-12-17T12:35:44 runtime: 1-02:16:12

User:

Proj: snic2019-1-12

Jobname: fesn_x

Flags: none



UPPSALA
UNIVERSITET

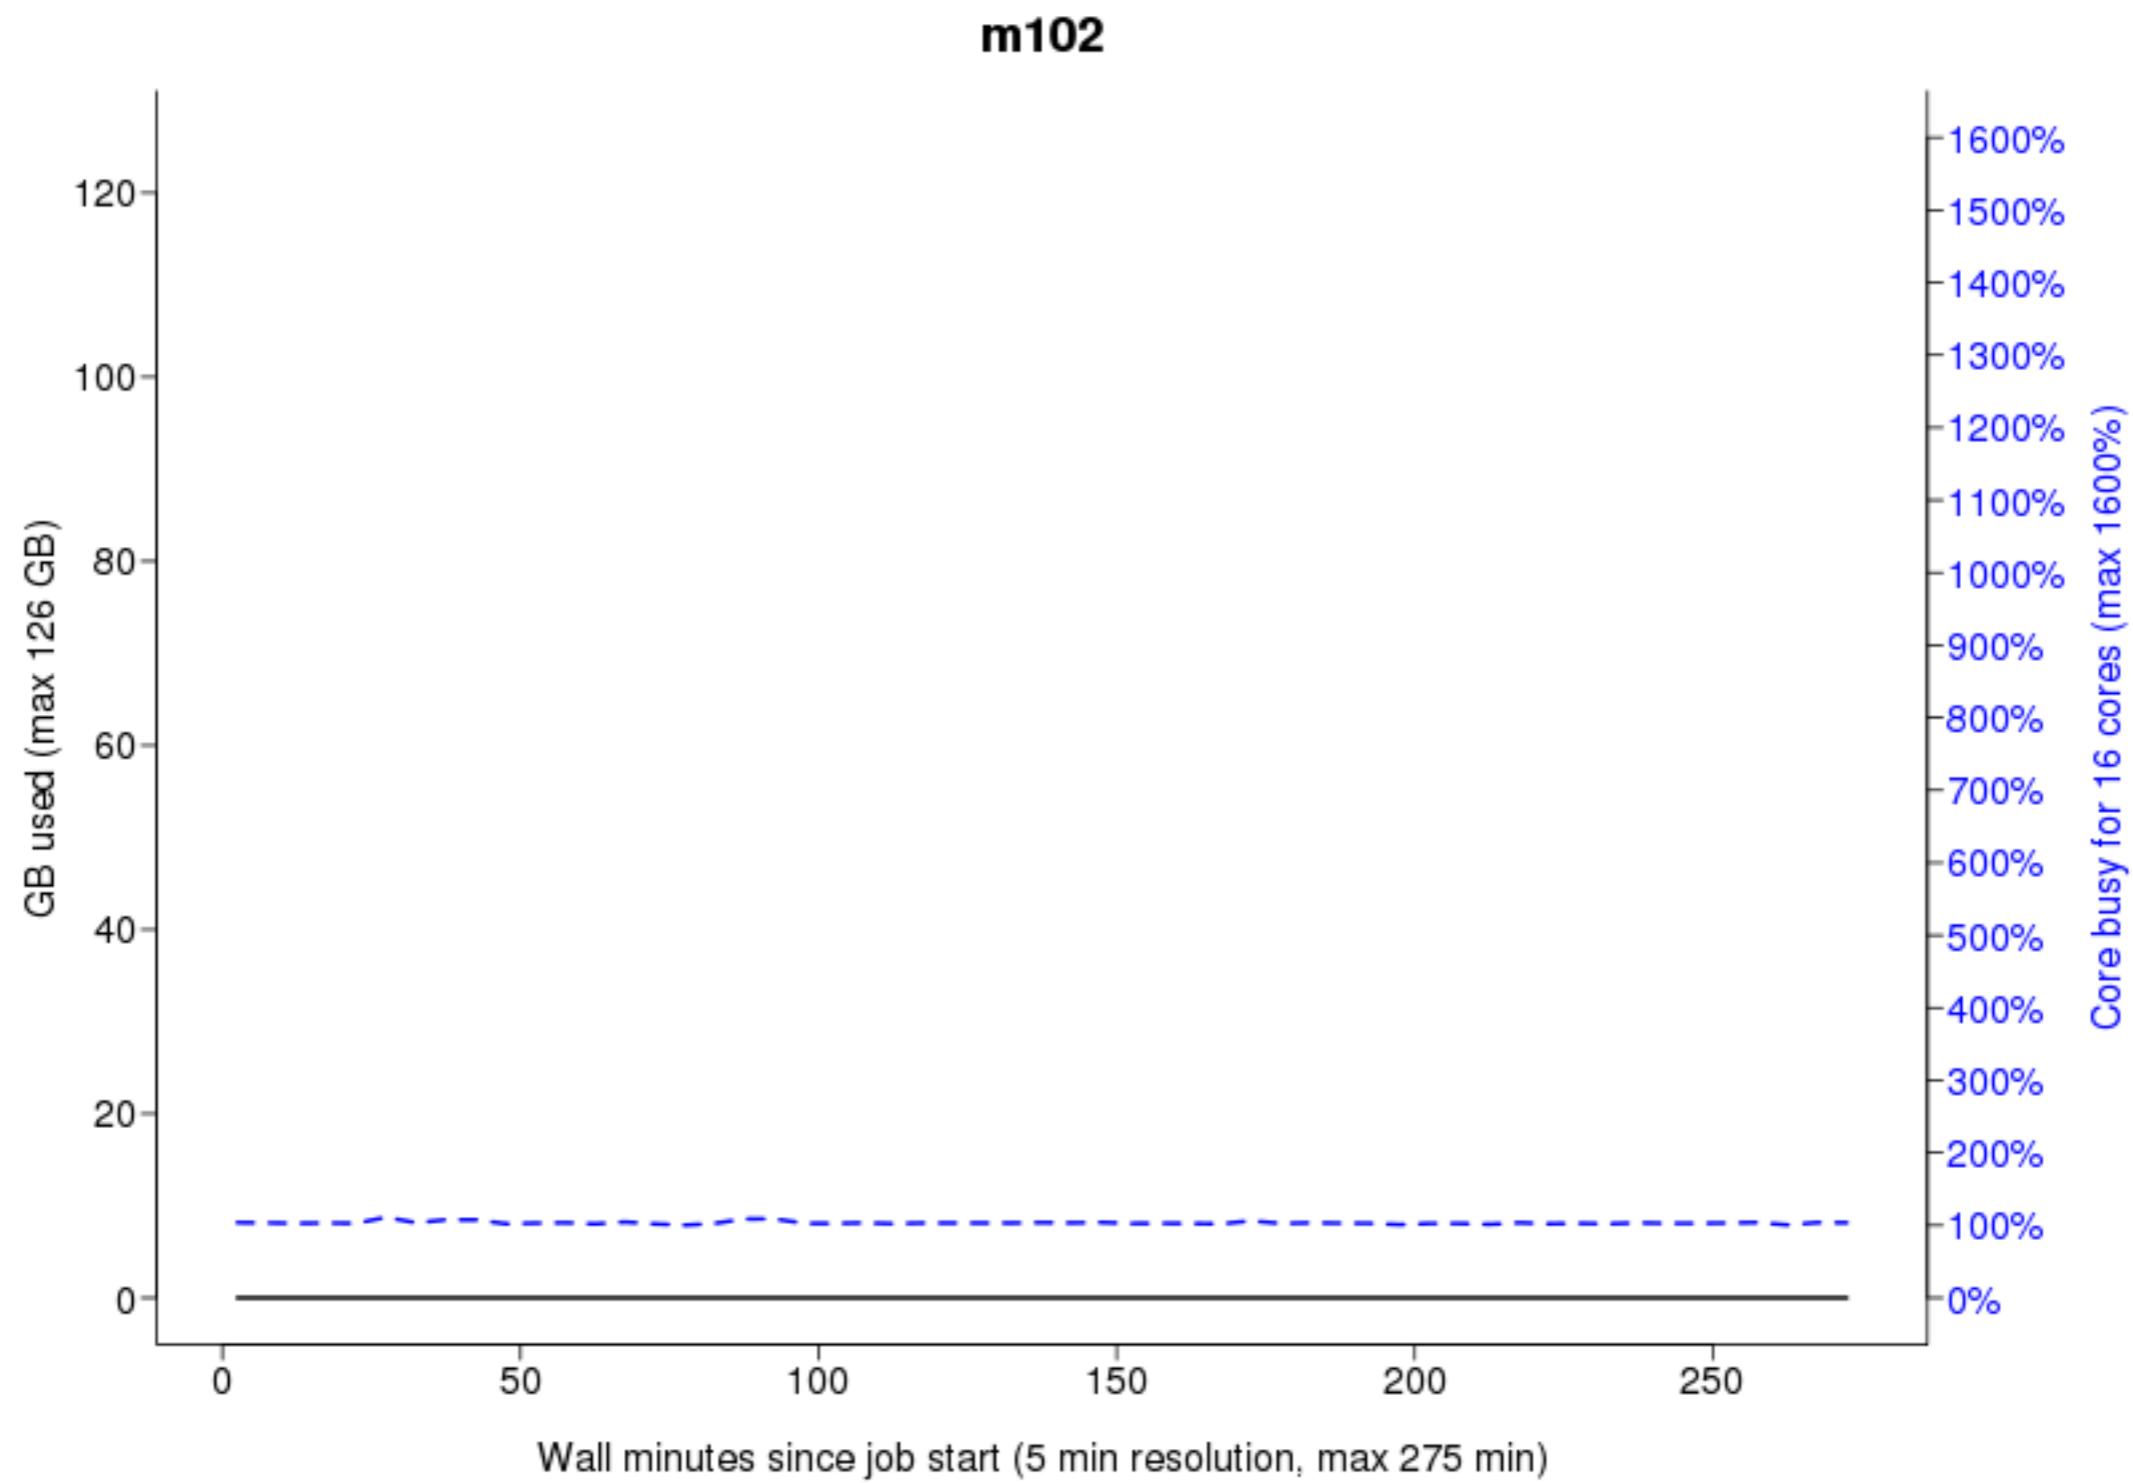
8804061 COMPLETED on milou end: 2016-10-10T20:30:26 runtime: 04:36:33

User:

Proj: b2015110

Jobname: nf-MergeBam_(2)

overbooked:12%, !!half_overbooked, !!severely_overbooked,
cores_overbooked:16:2, mem_overbooked:126:0, core_mem_overbooked:15.8:0



UPPSALA
UNIVERSITET

607031 OUT_OF_MEMORY on rackham end: 2018-08-25T07:29:39 runtime: 12:08:5

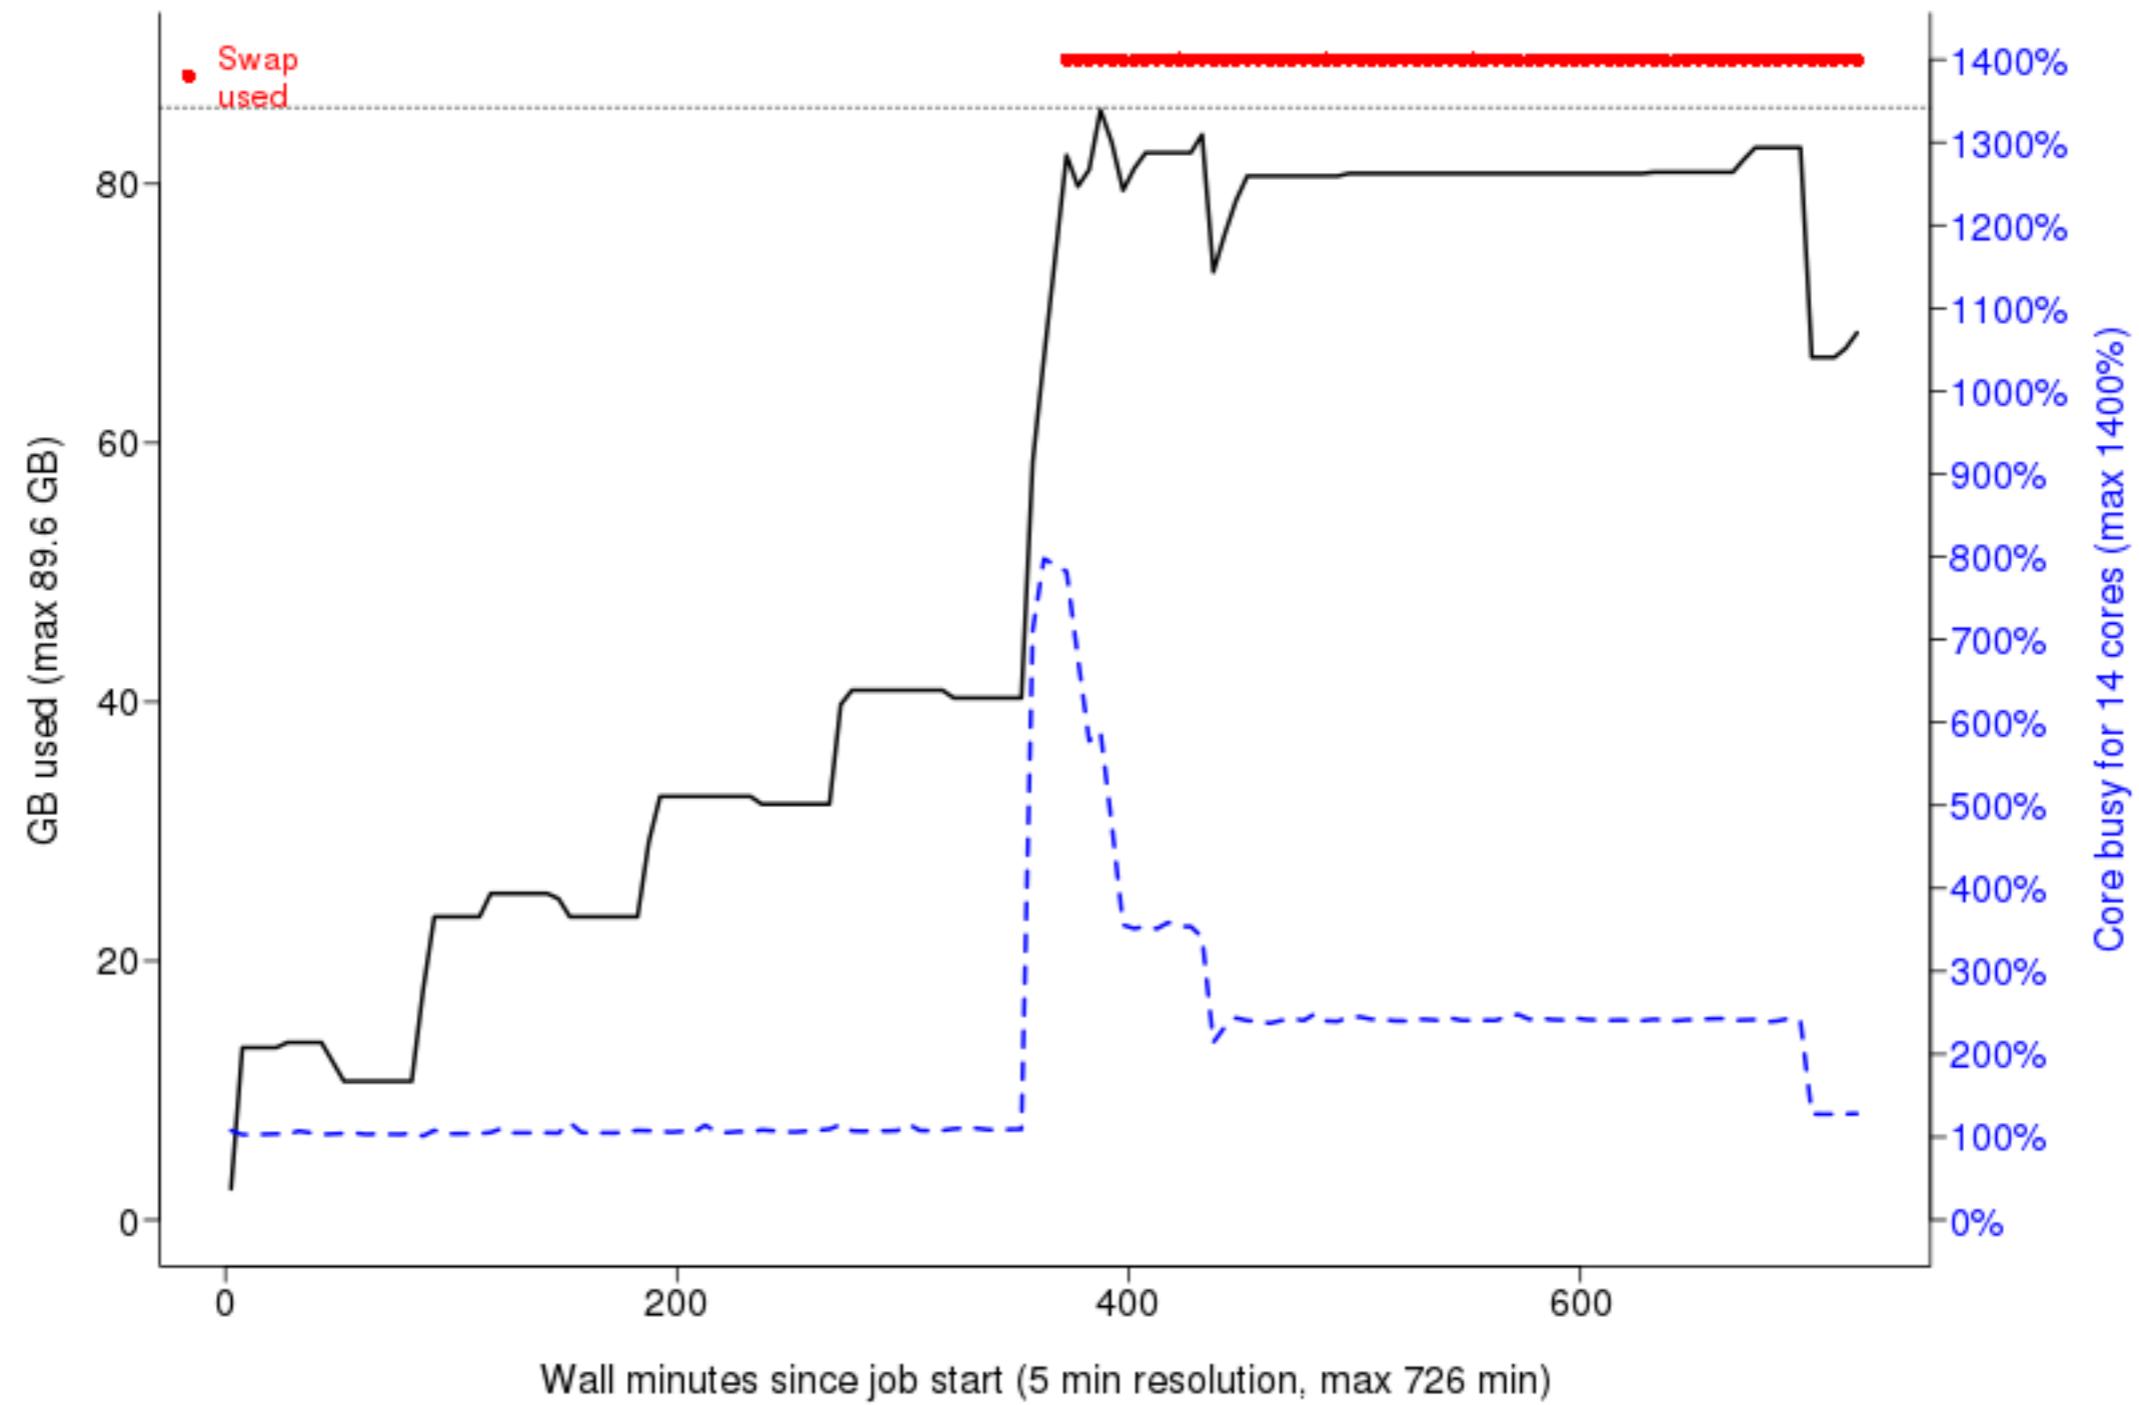
User:

Proj: snic2017-1-355

Jobname: pgd.apps

!!swap_used, cores_overbooked:14:11

r267



UPPSALA
UNIVERSITET



Different flavours of Slurm

Job script examples and workflows



Simple workflow

```
#!/bin/bash
#SBATCH -J jobname
#SBATCH -A uppmax2025-2-262
#SBATCH -p core
#SBATCH -n 10
#SBATCH -t 10:00:00
```

```
module load software/version
module load python/3.9.5
```

```
./my-script.sh
./another-script.sh
./myprogram.exe
```





Why do we need to load modules in the job script?

or in other words:

What does module load do?

Is it needed to specify the version of the software when loading the module? Why?



Job dependencies



- `sbatch jobscript.sh` submitted job with *jobid1*
- `sbatch anotherjobscript.sh` submitted job with *jobid2*
- `--dependency=afterok:jobid1:jobid2` job will only start running after the successful end of jobs *jobid1*:*jobid2*
- very handy for clearly defined workflows
- One may also use `--dependency=afternotok:jobid` in case you'd like to resubmit a failed job, OOM for example, to a node with a higher memory: `-C mem215GB` or `-C mem1TB`



I/O intensive jobs: \$SNIC_TMP

```
#!/bin/bash
#SBATCH -J jobname
#SBATCH -A uppmax2025-2-262
#SBATCH -p core
#SBATCH -n 1
#SBATCH -t 10:00:00

module load bioinfotools
module load bwa/0.7.17 samtools/1.14

export SRCDIR=$HOME/path-to-input

cp $SRCDIR/foo.pl $SRCDIR/bar.txt $SNIC_TMP/.
cd $SNIC_TMP

./foo.pl bar.txt

cp *.out $SRCDIR/path-to-output/.
```



OpenMP or multi-threaded job

```
#!/bin/bash
#SBATCH -A uppmax2025-2-262
#SBATCH --exclusive
#SBATCH -p node
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=20
#SBATCH -t 01:00:00
```

```
module load uppasd
export OMP_NUM_THREADS=20
```

```
sd > out.log
```



GPU nodes on Snowy



- Nodes with 1 Nvidia T4
- Available to everyone, priority to groups that paid for them

```
#SBATCH -M snowy
```

```
#SBATCH -C gpu
```

```
#SBATCH --gres=gpu:1
```

```
#SBATCH --gpus-per-node=1
```

- Further documentation:

https://slurm.schedmd.com/gres.html#Running_Jobs



Running on several nodes: MPI jobs

```
#!/bin/bash -l
#SBATCH -J rsptjob
#SBATCH --mail-type=FAIL
#SBATCH -A uppmax2025-2-262
#SBATCH -t 00-07:00:00
#SBATCH -p node
#SBATCH -N 4
### for jobs shorter than 15 min (max 4 nodes):
##SBATCH --qos=short
```

```
module load RSPT/2021-10-04
export RSPT_SCRATCH=$SNIC_TMP
```

srun -n 80 rspt

```
rm -f apts dmft_lock_file e_entropy efgArray.dat.0 efgData.out.0
energy_matrices eparm_last interstitialenergy jacob1 jacob2
locust.* out_last pot_last rspt_fft_wisdom.* runs.a symcof_new
```



Job arrays



- Submit many jobs at once with the same or similar parameters
- Use `$SLURM_ARRAY_TASK_ID` in the script in order to find the correct path

```
#!/bin/bash
#SBATCH -A uppmax2025-2-262
#SBATCH -p node
#SBATCH -N 2
#SBATCH -t 01:00:00
#SBATCH -J jobarray
#SBATCH --array=0-19
#SBATCH --mail-type=ALL,ARRAY_TASKS

# $SLURM_ARRAY_TASK_ID tells the script which iteration to run
echo $SLURM_ARRAY_TASK_ID

cd /path/to/my/directory/dir_${SLURM_ARRAY_TASK_ID}/

srun -n 40 my-program
env
```

- You may use `scontrol` to modify some of the job arrays.



Snakemake and Nextflow

- Conceptually similar, but with different flavours
- First define steps, each with an input, an output, and a command that transforms the input into output
- Then just ask for the desired output and the system will handle the rest





Hands-on #4: make it your own

- use 2 or 3 of the sample job scripts as a starting point for your own job script
- tweak them so that you run something closer to your research; or just feel free to experiment
- paste at least one of the examples in the HackMD
- great if you could add a comment what the job script is about





Feedback on Slurm

- what did you find useful?
- not so useful?
- what is most challenging while editing your job script / workflow?
- something that was not covered that you'd like to know about?
- please provide your feedback in the HackMD





Where to go from here?

Code documentation

NAISS training newsletter or <https://www.naiss.se/events/> - software-specific training events included

<https://coderefinery.org/workshops/upcoming/> (Git, Jupyter, code testing, writing code documentation, ...)

<https://nbis.se/training/events.html> (bio)

<https://enccs.se/events>

<https://supr.naiss.se/support/> or email support@uppmax.uu.se

